



1 Graph Traversal




(a) Recall that given a DFS tree, we can classify edges into one of four types:

- Tree edges are edges in the DFS tree,
- Back edges are edges (u, v) not in the DFS tree where v is the ancestor of u in the DFS tree
- Forward edges are edges (u, v) not in the DFS tree where u is the ancestor of v in the DFS tree
- Cross edges are edges (u, v) not in the DFS tree where u is not the ancestor of v , nor is v the ancestor of u .

For the directed graph above, perform DFS starting from vertex A, breaking ties alphabetically. As you go, label each node with its pre- and post-number, and mark each edge as Tree, Back, Forward or Cross.

(b) What are the strongly connected components of the above graph?




Red arrows indicate reversed edges

1. Run DFS on G (part a)
2. Run DFS starting from vertex with highest post number
3. Each time you have to restart represents a connected component

* Highest post number of DFS must lie in a source SCC in G (think how the only way to reach a vertex in a source SCC is to start at the source)

* A source SCC in G is a sink SCC in G^R

(c) Draw the DAG of the strongly connected components of the graph.



2 BFS Intro

In this problem we will consider the shortest path problem: Given a graph $G(V, E)$, find the length of the shortest path from s to every vertex v in V . For an unweighted graph, the length of a path is the number of edges in the path. We can do this using the *breadth-first search* (BFS) algorithm, which we will see again in lecture this week.

BFS can be implemented just like the depth-first search (DFS) algorithm, but using a queue instead of a stack. Below is pseudo-code for another implementation of BFS, which computes for each $i \in \{0, 1, \dots, |V| - 1\}$ the set of vertices distance i from s , denoted L_i .

```

1: Input: A graph  $G(V, E)$ , starting vertex  $s$ 
2: for all  $v \in V$  do
3:    $visited(v) = False$ 
4:  $visited(s) = True$ 
5:  $L_0 \leftarrow \{s\}$ 
6: for  $i$  from 0 to  $n - 1$  do
7:    $L_{i+1} = \{\}$ 
8:   for  $u \in L_i$  do
9:     for  $(u, v) \in E$  do
10:      if  $visited(v) = False$  then
11:         $L_{i+1}.add(v)$ 
12:       $visited(v) = True$ 
  
```

In other words, we start with $L_0 = \{s\}$, and then for each i , we set L_{i+1} to be all neighbors of vertices in L_i that we haven't already added to a previous L_i .

(a) Prove that BFS computes the correct value of L_i for all i (Hint: Use induction to show that for all i , L_i contains all vertices distance i from s , and only contains these vertices).

Base case: $i = 0$

True that $L_0 = \{s\}$, since we start at s , s is zero distance away from s .

Induction case:

Assume this holds for $i = k$. Will prove for $i = k+1$.

L_k contains all vertices distance k from s .

prove no vertex left out

Every vertex at distance $k+1$ is adjacent to a vertex in L_k .

No vertex at distance $k+2$ or more can be adjacent to a vertex in L_k .

prove there is no unwanted vertex

L_{k+1} set to all neighbors of vertices in L_k so L_{k+1} contains all vertices distance $k+1$ from s .

- (b) Show that just like DFS, the above algorithm runs in $O(m + n)$ time, where n is the number of nodes and m is the number of edges.

Init visited : $O(n)$

Each iteration: $O\left(\sum_{v \in I_i} \deg(v)\right)$

↳ eventually visit each edge twice (once per adjacent vertex)

Overall : $O(n + m)$

- (c) We might instead want to find the shortest *weighted* path from s to each vertex. That is, each edge has weight w_e , and the length of a path is now the sum of weights of edges in the path. The above algorithm works when all $w_e = 1$, but can easily fail if some $w_e \neq 1$.

Fill in the blank to get an algorithm computing the shortest paths when w_e are integers: We replace each edge e in G with w_e to get a new graph G' , then run BFS on G' starting from s . Justify your answer.

$A \xrightarrow{3} B$ becomes $A \xrightarrow{1} \circ \xrightarrow{1} \circ \xrightarrow{1} B$

- now each edge has weight 1

- only one neighbor to visit after each dummy node, so path is still correct

- (d) What is the runtime of this algorithm as a function of the weights w_e ? How many bits does it take to write down all w_e ? Is this algorithm's runtime a polynomial in the input size?

of edges: $\sum_{e \in E} w_e$

Runtime : $O\left(\sum_{e \in E} w_e + n\right) = O\left(\sum_{e \in E} w_e\right)$

bits to represent the number N : $\log N$

bits to write down all w_e : $\sum_{e \in E} \log w_e$

Not polynomial, runtime is exponential in input size.


For runtime to be polynomial it should run in $O\left(\left(\sum_{e \in E} \log w_e\right)^c\right)$

for some constant $c > 0$.


3 Dijkstra's Algorithm Fails on Negative Edges

Draw a graph with five vertices or fewer, and indicate the source where Dijkstra's algorithm will be started from.

- (a) Draw a graph with no negative cycles for which Dijkstra's algorithm produces the wrong answer.



- (b) Draw a graph with at least two negative weight edges for which Dijkstra's algorithm produces the correct answer.




4 Waypoint


You are given a strongly connected directed graph $G = (V, E)$ with positive edge weights, and there is a special node $v_0 \in V$. Give an efficient algorithm that computes, for all node pairs s, t , the length of the shortest path from s to t that passes through v_0 . Your algorithm should take $O(|V|^2 + |E| \log |V|)$ time.

Observation 1: shortest path $s \rightarrow v_0 \rightarrow t$ is broken down into two shortest paths $s \rightarrow v_0$ and $v_0 \rightarrow t$

Observation 2:



Each of these paths can be found with one iteration of Dijkstra's from v_0 on the graph:



Algorithm:

1. Compute shortest paths $v_0 \rightarrow t$ by running Dijkstra's on v_0 in G
2. " " " $s \rightarrow v_0$ " " " " v_0 in G^R
3. Iterate over results of 1 and 2 to piece together all shortest paths between all pairs

$O(|E| \log |V|)$
 $O(|V|^2)$

5 Dijkstra Tiebreaking

We are given a directed graph G with positive weights on its edges. We wish to find a shortest path from s to t , and, among all shortest paths, we want the one in which the longest edge is as short as possible. How would you modify Dijkstra's algorithm to this end? Just a description of your modification is needed.

← tie break based on longest edge

(If there are multiple shortest paths where the longest edge is as short as possible, outputting any of them is fine).

ex: $A \xrightarrow{1} B \xrightarrow{1} C \xrightarrow{1} D$ * we choose this as THE shortest path
 $A \xrightarrow{1} E \xrightarrow{2} D$

Modify Dijkstra's to keep a map $l(v) \rightarrow$ longest edge in shortest path to v

Init $l(s) = 0$ and $l(v) = \infty$ for all $v \in V \setminus \{s\}$

When considering the edge $u \rightarrow v$

if $\text{dist}(u) + w(u, v) < \text{dist}(v)$:

$\text{edgeTo}(v) = u$

$l(v) = \max(l(u), w(u, v))$

if $\text{dist}(u) + w(u, v) = \text{dist}(v)$ and $l(v) > \max(l(u), w(u, v))$:

$\text{edgeTo}(v) = u$

$l(v) = \max(l(u), w(u, v))$